

Validation platform: market, business plan, roadmap and backlog

Sep 30, 2026 · @Mihai

Market

Nobody sells a lightweight, self-serve tool for validating a requirement list with stakeholders. The need is served from two sides that leave a gap in the middle: enterprise sourcing and requirements suites priced from \$8,000 per project upward, and \$19 to \$129 per month feedback boards built for customer feature voting, not expert validation. In between, most teams use Excel plus a survey form.

Product	What it does	Pricing (as found)	Where it falls short for this use
Olive	AI-native vendor evaluation: requirements set, stakeholder surveys, weighting, vendor scoring, final report. Closest competitor for platform-selection validation.	Free research tier without collaboration; \$8,000 per single project; platform from \$24,000 per year, billed annually	Built around RFPs and vendor scoring. Overkill and overpriced for a PM validating a roadmap or a catalogue; no self-serve paid entry below \$8,000
Jama Connect	Requirements management with a Review Center: stakeholders review, comment, approve or reject, sign off	Quote-based, per user; AWS Marketplace lists it as private offer only	Regulated-industry engineering tool. Weeks to configure, admin-driven, stakeholders need licences

Product	What it does	Pricing (as found)	Where it falls short for this use
ProductLift	Feedback boards, public roadmap, changelog; RICE, ICE, MoSCoW scoring; anonymous voting; EU-hosted	Starter \$19, Pro \$49, Business \$129 per month, unlimited voters	Customer upvoting model. No per-item agree-or-push-back with reasons, no proposed-priority validation, no respondent attributes or perspective filtering
Canny, Frill, Featurebase	Same category as ProductLift	Canny free to 25 users then from \$19 per month; Frill \$25 to \$349 per month	Same gap: votes, not structured validation
Typeform, Microsoft Forms, Google Forms	Generic surveys	Free to about \$50 per month	One question per item is unreadable at 80 items; no grouping, no proposed values, no per-item reasons, no results view a PM can act on
Miro, Slido, Mentimeter	Workshop voting and dot-voting templates	Free to about \$20 per user per month	Synchronous only; no resume, no structured record, no register of disagreements

The gap this product fills: a PM imports a list, gets a readable instrument in minutes, sends one link, and gets back agreement, disagreement with reasons and gaps, at \$50 to \$200 per month rather than \$8,000 per project.

Three things the search confirmed. First, the mechanics are public: MoSCoW voting, stakeholder surveys and consensus dashboards appear in ProductLift, Olive, Miro templates and procurement guidance. Second, AI-drafted requirements are already a selling point at Olive, so the Claude feature is expected, not novel; what is not on the market is AI that reads the responses and tells the PM what to do next. Third, Olive's own pricing page names the stack it replaces (spreadsheets, survey tools, Word reports, email threads), which is the same stack this product replaces at one fortieth of the price.

Sources opened: [Olive pricing](#), [Jama reviews help page](#), [ProductLift MoSCoW page](#), [ProductLift feature voting comparison](#). Canny, Frill, survey and workshop tool prices are from ProductLift’s comparison pages and are approximate.

Business plan

A hobby build that can become a EUR 3,000 to 5,000 per month business inside 18 months, sold to consultancies and in-house product teams, run by Alerty S.R.L., with zero fixed cost until the first paying customer.

Product in one line. Import a requirement list from anywhere, get a readable validation instrument in minutes, send one link, get back agreement, disagreement with reasons and gaps, plus AI-written actions.

Customers, in order of attack.

- 1. Boutique consultancies and agencies (5 to 50 people) that run discovery and platform selection for clients. They validate lists every engagement, bill by the hour, and can charge the tool to the client. Buyer: delivery lead or partner. Olive proves this segment pays, at \$8,000 per project.
- 2. In-house product and IT teams in mid-size and large companies validating roadmaps, catalogues and platform requirements with internal SMEs. Buyer: head of product, PMO lead, IT sourcing lead. Slower to close; single-tenant and self-hosted options matter here.
- 3. Individual PMs. Free tier only; they are distribution, not revenue.

Constraint from the 3Pillar contract. No selling to 3Pillar clients met through the engagement until two years after the contract ends, and no use of the engagement’s files or their concept in the design. Segment 1 excludes 3Pillar itself unless they license it; segment 2 excludes Informa.

Pricing.

Plan	Price	Includes	For
Free	EUR 0	1 active project, 25 responses, 3 AI runs per month, branding shown	Individual PMs, trials

Plan	Price	Includes	For
Pro	EUR 49 per month or EUR 490 per year	Unlimited projects, 500 responses per month, AI shaping and insights, custom branding, exports	Freelancers, one PM
Team	EUR 149 per month or EUR 1,490 per year	5 seats, 2,500 responses per month, personal invites and reminders, Jira and Notion connectors, priority support	Consultancies, product teams
Enterprise	From EUR 6,000 per year	Single-tenant or self-hosted, SSO, audit log, DPA, unlimited responses, bring-your-own AI key	Large companies

Anchor: a consultancy pays EUR 1,490 per year for what Olive charges \$8,000 per project. AI usage is metered inside the response allowance so a heavy user cannot run up an unbounded Anthropic bill.

Costs, monthly, at three stages.

Item	Hobby (0 to 10 users)	First customers (10 to 100)	Business (100 to 500)
Hosting (Vercel Pro, Postgres, storage)	EUR 0 to 25	EUR 60 to 120	EUR 250 to 500
Anthropic API	EUR 5 to 20	EUR 50 to 150	EUR 300 to 800
Email, monitoring, analytics	EUR 0	EUR 30	EUR 80
Domain, legal templates, accounting for the SRL	EUR 15	EUR 40	EUR 100
Claude subscription for the build	already paid	already paid	already paid
Total	under EUR 60	under EUR 350	under EUR 1,500

Figures are approximate and assume EU-region Postgres on a managed provider and Anthropic usage capped per plan.

Revenue path. 20 Pro and 5 Team customers is EUR 1,725 per month. 50 Pro, 15 Team and 2 Enterprise is about EUR 5,700 per month. The second point is 18 months out if the first paying customer arrives within six months of launch. Hobby pace makes every date slip; a fixed weekly budget of hours is the only defence.

Distribution. Free tier and a sample project a visitor can click through without signing up; an offline export as a lead magnet for organisations that cannot use hosted tools; content on validation practice (MoSCoW done badly, why surveys fail at 80 items); LinkedIn, PM communities, Product Hunt at R2; partner programme for consultancies at R3.

Risks.

Risk	Likelihood	Mitigation
3Pillar contract dispute	Low, rises with visibility and revenue	Clean design record, no engagement files, no 3Pillar clients; consider asking for a waiver once the product exists and is visibly different
Nobody pays; teams stay on Excel and Forms	Medium	Free tier proves usage before pricing; the insight layer, not the form, is the paid value
Olive or ProductLift adds this in a quarter	Medium	Speed and focus; write-back to Jira and live mode are cheap for you and off-strategy for both
AI cost or quality embarrasses the product	Medium	Golden set evals, PM accepts every rewrite, metered usage, AI-off mode
Hobby stalls	High	Weekly hour budget, public roadmap, one real user by week 12
GDPR and processor obligations	Certain	DPA template, EU hosting, subprocessor list, deletion per workspace from R1

Build setup

You own accounts, secrets, decisions and acceptance; Claude Code owns the code, tests and documentation, run as a lead with subagents from the day 0 brief (agents, CLAUDE.md, tester never fixes, reviewer with fresh context). Nothing from the engagement enters the repository or any Claude session that designs the instrument or the dashboard.

How a working session runs. You open Claude Code on the repo, name the stories in scope, approve the split, and come back to a diff, a screenshot and a test command per story. You click through on a laptop and a phone and accept or reject with a plain note. Two to three hours of your time per session; budget two sessions a week.

Stack. Chosen for self-hosting later without a rewrite: nothing on the critical path is tied to one vendor.

Layer	Choice	Why
App	Next.js 15, TypeScript, Tailwind, shadcn/ui components	One codebase for PM side, respondent side, API routes; portable to a Docker container
Database	Postgres via Drizzle ORM, hosted on Neon (EU region) in SaaS, plain Postgres in self-host	Plain SQL, migrations in the repo, no vendor features
Auth	better-auth (magic link, Google, Microsoft, later SAML)	Library in the app, not a hosted identity provider
Realtime	Server-sent events from the app	Works in any host, no vendor channel
Files	S3-compatible storage: Cloudflare R2 (EU jurisdiction) in SaaS, MinIO in self-host	Same client either way
Email	Resend, with Postmark as the alternative if EU processing becomes a customer requirement	Transactional only
AI	Anthropic API from server routes; enterprise bring-your-own key, Bedrock and Vertex in R3	Key never in the browser

Layer	Choice	Why
Hosting	Vercel, region fra1; Docker Compose for self-host in R3	Fast to ship; container path kept open
Payments	Stripe with Stripe Tax	EU VAT and OSS handled
Monitoring	Sentry (EU data region), Plausible	Errors, usage, no cookies
Tests	Vitest, Playwright	Unit for logic, browser for every user flow

Tools on your machine. Claude Code (desktop app), Git, Node 22, Docker Desktop (needed from R1 for a local Postgres and later for the self-host package), a browser with the Claude in Chrome extension or the Playwright MCP so Claude can see what it built.

Accounts and secrets, by phase. You create every account and paste keys into `.env.local` and the Vercel project; Claude never sees or commits a key.

Phase	Account	What to prepare	Cost
Setup	GitHub	Organisation under Alerty, private repo, branch protection on main	Free
Setup	Domain registrar (Cloudflare Registrar or Namecheap)	The product domain plus a mail subdomain for sending	EUR 10 to 40 per year
Setup	Anthropic Console	API key, monthly spend limit set to EUR 50	Usage
Setup	Neon	Project in EU (Frankfurt), one branch for dev, one for prod	Free to EUR 19 per month
Setup	Vercel	Team under Alerty, project linked to the repo, region fra1, env vars	Free to EUR 20 per month
Setup	Google Cloud Console	OAuth client for Google sign-in, consent screen	Free
Setup	Microsoft Entra	App registration for Microsoft sign-in	Free

Phase	Account	What to prepare	Cost
R1	Resend	Verified sending domain (SPF, DKIM, DMARC records at the registrar)	Free to EUR 20 per month
R1	Cloudflare R2	Bucket for uploads and exports, EU jurisdiction flag	Free tier
R1	Sentry	Project created with EU region selected at creation	Free tier
R1	Plausible	Site added, script in layout	EUR 9 per month
R1	Legal templates	Privacy policy, terms, DPA with subprocessor list; use a generator then have a Romanian lawyer read them once	EUR 0 to 500 one-off
R2	Atlassian developer	Jira OAuth app for import and write-back	Free
R2	Notion, Trello developer	Integration tokens	Free
R3	Stripe	Alerty verified, Stripe Tax on, products and prices created	Fees per transaction
R3	Docker Hub or GitHub Container Registry	Image publishing for self-host	Free
R3	An SSO test tenant (Okta developer or Entra)	SAML and OIDC testing	Free


Repository skeleton on day 0. `CLAUDE.md`, `WRITING.md` (your AI-tells guide), `SECURITY.md`, `INTERFACES.md`, `MISTAKES.md`, `.claude/agents/` with the nine agents, `stories/` one file per story, `evals/` with the golden set, `docs/design-notes/` dated, `docs/decisions/` one file per architecture decision.

Roadmap

Six phases from a two-week setup to a first payment in summer 2027, at two sessions a week. Each phase ends at a gate you can check in one sentence; a missed gate means the next phase waits, it does not start anyway.

Six phases to a first payment by August 2027

Hobby pace: two sessions a week, about eight hours. Diamonds are gates.



roadmap · 6 phases, 6 gates, Oct 2026 to Aug 2027

R1 is the long band because it is the product; everything before it is preparation and everything after it depends on what its first users say. Design sits before R1 so the instrument and the dashboard are drawn from the problem and public tools, with dated notes in the repo, before any code that renders them exists.

What each phase produces:

1. Setup: accounts and secrets from the Build setup table, the repository skeleton, CLAUDE.md and the nine agents, the first story map, a name and a domain.
2. Design: the respondent instrument and the dashboard on paper and as clickable prototypes, the design system (type, colour, spacing, components), the landing page copy, the golden set of ten messy specs for the AI evals.
3. R1 build: the epics below, in order. Ends with the app live at the domain on the free plan and one organisation you do not work for running a real validation.
4. First users: no new features. Fix what hurt, rewrite the landing page from what respondents and PMs said, publish the offline lead magnet, post the first content.
5. R2 loop: Jira and Notion import and write-back, set versioning with diffs between rounds, respondent questions answered from the source, live workshop mode, team roles.
6. R3 revenue: Stripe and plans, single-tenant hosting, the Docker self-host package with bring-your-own AI key, SSO, audit log, DPA flow.

Epics and stories for R1

Twelve epics, built in this order; each story is one Claude Code session or less and ends with a test and a screenshot. Acceptance criteria are the check you do in the evening. Stories in `stories/` get the same text plus the technical notes the builder needs.

E1 Foundation. Repo, CI, database, deploy pipeline.

Story	Acceptance
Repo skeleton with CLAUDE.md, agents, companion files, lint, format, commit hooks	<code>npm run lint</code> and <code>npm test</code> pass on an empty app; CI runs both on every push
Postgres schema v1 for the object model (workspace, project, requirement set, item, instrument, invite, response, answer, insight) with migrations	Migration applies to an empty database and rolls back; schema documented in <code>docs/schema.md</code>
Deploy to Vercel fra1 from main with preview deploys per branch	A hello page at the domain over HTTPS with security headers scored A on securityheaders.com
Local development with Docker Postgres and seed data	<code>docker compose up</code> plus <code>npm run dev</code> gives a running app with one seeded workspace

E2 Accounts and workspaces.

Story	Acceptance
Sign in with magic link	Email arrives within 30 seconds, link signs in once, expires after 15 minutes
Sign in with Google and Microsoft	Both providers create or match the same user by email
Create a workspace on first sign-in, name it, switch between workspaces	New user lands in a named workspace; a second workspace appears in the switcher
Invite a member by email, owner and member roles	Member can create projects, cannot delete the workspace or change billing
Workspace settings: name, logo, accent colour, AI budget shown	Logo and colour appear on every instrument created after saving

E3 Import.

Story	Acceptance
Upload xlsx or csv, detect header row, preview first ten rows	Files up to 5 MB and 2,000 rows import; a file without a header row asks which row is the header
Column mapping to item text, area, proposed value, reference, and up to five custom fields	Mapping remembered per workspace for files with the same headers
Import log and set versioning	Re-importing creates version 2; version 1 stays readable and attached to its instruments
Import validation report	Empty items, duplicates by exact text and rows over 1,000 characters listed before commit

E4 AI shaping.

Story	Acceptance
Server route to Anthropic with per-workspace budget, request logging, timeout	A workspace over budget gets a clear message; no key in any client bundle
Group items into areas and order the areas with a one-line rationale each	Every item lands in exactly one area; the PM can drag items between areas
Reader version per item with the original kept and shown on tap	No rewrite is applied without an explicit accept; accept all and reject all exist
Duplicate and ambiguity flags	Flags link to the items involved; dismissing a flag is remembered
Golden set of ten messy specs in <code>evals/</code> with expected areas and item counts; eval runner reports found, missed, invented, meaning-changed	Runner executes in CI on any prompt change and fails on a new invented item

E5 Instrument builder.

Story	Acceptance
Create an instrument from a set version; intro text; respondent fields with mandatory flag, text or dropdown	Preview refuses to start until mandatory fields are filled
Scoring method templates: MoSCoW with proposed value, 1 to 5 fit, keep change drop, custom labels	Switching method re-renders every item; answers are not kept across a method change
Layout templates: one item per screen, one area per screen, single long page	All three pass the phone check at 375 px wide
Perspective filter: tag items, respondent picks perspectives, sees tagged items only	Untagged items are shown to everyone; the dashboard shows coverage per perspective
Closing questions: free text and a mandatory confidence scale	Instrument cannot be published without at least the confidence question
Preview as respondent on desktop and phone frames	Preview responses are never stored

E6 Sharing.

Story	Acceptance
Publish with a public link carrying a 128-bit token; open and close dates; optional passcode	A closed instrument shows a closed page, not an error; a wrong passcode is rate-limited
Personal invites: paste emails, send from the platform, one link per person	Each personal link resumes that person's response on any device
Reminders to invitees who have not submitted, at most one every three days, manual trigger	Reminder count and last sent date shown per invitee
Kill switch: revoke the public link or one personal link	Revoked link shows a revoked page within one minute

E7 Respondent experience.

Story	Acceptance
Open a link, fill mandatory fields, start	Works without an account on iOS Safari, Android Chrome, desktop Chrome, Edge, Firefox
Rate items in the chosen method; a score under agreement opens a reason box; unclear opens a question box; optional comment on any item	Reason boxes are mandatory before the item counts as answered
Autosave every answer within one second; resume by personal link or by browser token on a public link	Closing the tab mid-way and reopening restores every answer
Navigation by area with per-area progress; overall progress	An item never rated is distinguishable from one skipped on purpose
Missing-item submission as a structured item (text, area, suggested value)	Appears in the dashboard's missing register with the respondent's name
Submission summary and submit; edits allowed until close date	Submission timestamp stored in UTC; a second submit updates, never duplicates
Accessibility: keyboard only, screen reader labels, contrast AA	Playwright axe scan reports zero serious violations

E8 Dashboard.

Story	Acceptance
Response tracker: who, status, progress, submitted date, source	Filters by respondent field and by perspective
Agreement per item and per area with distribution of answers	Numbers reconcile with a CSV export to the row
Disagreement register: item, respondent, their value, their reason	Sortable by item and by respondent; click-through to the item detail
Unclear register and missing-item register	Each row links to the item or to the suggested new item

Story	Acceptance
Item detail: every respondent's answer and comment on one item	Loads under 500 ms with 100 responses
Conflict view: where respondent groups disagree, by any respondent field	Shows the four largest gaps by default
Live updates while the instrument is open	A new answer appears on the dashboard within five seconds without reload
Sample project on first login, watermarked, deletable	Watermark cannot be removed except by deleting the sample

E9 Insights.

Story	Acceptance
Generate actions from responses: items to rewrite, group conflicts, follow-ups, low-coverage areas	Every action cites at least one response; an action without a citation is dropped before display
Mark an action done or dismissed	Regeneration does not resurrect a dismissed action
Insight cost shown per run against the workspace budget	A run that would exceed budget is refused with the shortfall shown

E10 Exports.

Story	Acceptance
CSV of responses and of items with aggregates	Opens in Excel with correct UTF-8 and dates
JSON export of the full project	Re-importable into another workspace
PDF summary: headline numbers, per-area agreement, registers	Renders under 10 seconds for 200 items and 50 responses

E11 Trust and compliance.

Story	Acceptance
Rate limiting on public links and auth routes	100 requests per minute per IP on respondent routes; burst on auth blocked after 5 attempts
Data export and deletion per workspace, self-serve	Deletion removes all rows and files within 24 hours and emails a confirmation
Privacy policy, terms, DPA and subprocessor list pages; privacy notice on every instrument	Pages reviewed once by a lawyer; version and date shown
Backups nightly with a tested restore	A restore into a fresh database is performed and documented once before launch
Sentry and Plausible wired, no personal data in events	Sentry scrubs emails; Plausible runs without a cookie banner

E12 Landing and onboarding.

Story	Acceptance
Landing page: problem, product, three-step how it works, sample instrument you can click, pricing, sign up	Lighthouse performance and accessibility over 90 on mobile
Quickstart in the app: import a sample file, publish, open the link on your phone	A new user reaches a published instrument in under five minutes without help
Transactional emails: magic link, invite, reminder, submission receipt to the PM	Pass the WRITING.md scan; render in Gmail, Outlook web and Apple Mail

Visual design

No designer for R1. Claude can produce a site and an app at the level of trello.com in layout, typography, colour, components, responsiveness and motion, as working code rather than mockups. Two things it cannot produce and Trello has: a consistent hand-drawn illustration set, and a distinctive logo. Both can be bought for under EUR 1,500 or replaced by a design language that does not need them.

What trello.com is made of. A mega-menu navigation with grouped feature links, a hero with an email sign-up and a product video, a slider of illustrated product screens (Inbox, Boards, Planner), feature sections with spot illustrations in SVG, a testimonial carousel with survey statistics and customer logos, a plan comparison strip, a learn-and-connect grid, and a footer with a 21-language switcher. Underneath is the Atlassian Design System: one type scale, one colour system, one component set, used on every page.

What Claude covers, with the current toolchain.

Element	Claude can	How
Page structure and layouts	Yes, at that level	Next.js pages, CSS grid and flex, responsive to 320 px
Typography system	Yes	Google Fonts pairing, a type scale, line lengths and spacing tokens
Colour system, light and dark	Yes	Tokens in CSS variables, contrast checked to WCAG AA
Component library	Yes	shadcn/ui on Radix, restyled to the product's tokens; every state (hover, focus, disabled, loading, empty, error)
Product screens as hero art	Yes	Real screens of the app, framed and lightly styled, the way Trello's sliders show its boards
Motion	Yes	CSS transitions and Framer Motion for reveals, sliders, progress
Icons	Yes	Lucide (open licence), consistent stroke weight
Diagrams and abstract art	Yes	Hand-written SVG: flows, shapes, gradients, patterns
Emails	Yes	React Email templates tested in the main clients
Mega-menu, carousel, language switcher, pricing table	Yes	Standard components, accessible by keyboard

Element	Claude can	How
Character or scene illustrations in one hand-drawn style	No	Claude's SVG art reads as diagrammatic; a set of eight matching scenes is illustrator work
A distinctive logo mark	Partly	A wordmark in a chosen typeface, yes; an original mark that survives at 16 px and on a badge, no
Photography	No	Stock or none
Seeing its own work	Only with a browser	In Claude Code with the Playwright MCP or the Chrome extension, Claude screenshots every screen at three widths and fixes what it sees; without that loop, quality drops

Two ways to close the gaps.

1. Buy the two missing pieces. An illustrator on a freelance marketplace for eight spot illustrations in one style, EUR 300 to 1,000; a logo designer for a mark plus wordmark, EUR 200 to 800. Claude writes the brief, the palette and the reference board for both.
2. Choose a design language that does not depend on them. Product-forward and typography-led: real app screens as the art, strong type, a restrained two-colour palette with one accent, abstract SVG shapes for texture. This is the language of most current B2B tools and it is what a PM buyer reads as serious. Trello's playful illustration style is a consumer-productivity signal you do not need.

Recommendation: option 2 for R1, with a wordmark; buy the logo mark in First users if the name sticks; buy illustrations only if the landing page tests flat without them.

Process. In Design phase Claude writes `docs/design-system.md` (type, colour, spacing, components, motion, voice) and builds a styleguide page in the app. Every screen is then built from it and reviewed by you on a laptop and a phone before it is accepted. The Trello site is a quality reference, not a template: the product gets its own layout and identity, and nothing is copied from Atlassian.